

PERFORMING KNN ON AMAZON FINE FOOD REVIEWS

August 3, 2018

1 KNN On Amazon Fine Food Reviews:

In []: => Predictions:

The purpose of this analysis **is** to make up a prediction model where we will be able to predict whether a recommendation **is** positive **or** negative. In this analysis, we will **not** focus on the Score, but only the positive/negative sentiment of the recommendation. To do so, we will work on Amazon's recommendation dataset, we will build a Term-doc incidence matrix using term frequency **and** inverse document frequency ponderation.

When the data **is** ready, we will load it into predictive algorithms, mainly naïve Bayesian and regression. In the end, we hope to find a **"best"** model **for** predicting there commendation's **sentiment**.

=> Loading the data:

In order to load the data, we will use the SQLITE dataset where we will only fetch the Score **and** the recommendation summary.

2. As we only want to get the **global** sentiment of the recommendations (positive **or** negative), we will purposefully ignore **all** Scores equal to

3. If the score **id** above 3, then the recommendation will be **set** to **"postive"**. Otherwise, it will be **set** to **"negative"**.

The data will be split into an training **set and** a test **set with** a test **set** ratio of 0.2

In []: => Attribute information:

1. Id
2. Product **id**
3. User **id**
4. Profile name
5. Helpful numerator
6. helpful denominator

7.Reviews=Positive (4 or 5) and Negative (1 or 2)
8.Time
9.Summary
10.Text
Objective:

1. Applied K-Nearest Neighbour on Different Featurization of Data viz. BOW, tfidf, Avg-Word2Vec(using Word2Vec model pretrained on Google News) and tf-idf-Word2Vec.
2. Used both brute & kd-tree implementation of KNN
- 3.Perform train and test split using time based slicing.
- 4.perform 10 fold cross validation to find optimal "k" in knn
- 5.Perform test accuracy on
 - A.BOW
 - B.Tfidf
 - C.Avg word2vec
 - D.Tfidf W2v

```
In [51]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")
import re
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
from time import time
import random
import gensim
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation
```

Using the table to read data

```
df=pd.read_csv('Reviews.csv')
df.head()
```

```
Out [51]:
```

	Id	ProductId	UserId	ProfileName	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	
2	3	B000LQOCHO	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham	"M. Wassir"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	1	1	5	1303862400	
1	0	0	1	1346976000	
2	1	1	4	1219017600	
3	3	3	2	1307923200	
4	0	0	5	1350777600	

	Summary	Text
0	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	"Delight" says it all	This is a confection that has been around a fe...
3	Cough Medicine	If you are looking for the secret ingredient i...
4	Great taffy	Great taffy at a great price. There was a wid...

```
In [52]: filtered_data=pd.read_csv('Reviews.csv')
filtered_data.head()
```

```
Out [52]:
```

	Id	ProductId	UserId	ProfileName	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	
2	3	B000LQOCHO	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham	"M. Wassir"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	1	1	5	1303862400	
1	0	0	1	1346976000	
2	1	1	4	1219017600	
3	3	3	2	1307923200	
4	0	0	5	1350777600	

	Summary	Text
0	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	"Delight" says it all	This is a confection that has been around a fe...
3	Cough Medicine	If you are looking for the secret ingredient i...
4	Great taffy	Great taffy at a great price. There was a wid...

```
In [53]: def partition(x):
         if x<3:
             return 'negative'
         return 'positive'
```

```
In [54]: actualScore=filtered_data['Score']
         positiveNegative=actualScore.map(partition)
         filtered_data['Score']=positiveNegative
         filtered_data.shape
         filtered_data.head()
```

```
Out [54]:
```

	Id	ProductId	UserId	ProfileName	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	
2	3	B000LQOCHO	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham	"M. Wassir"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	1	1	positive	1303862400	
1	0	0	negative	1346976000	
2	1	1	positive	1219017600	
3	3	3	negative	1307923200	
4	0	0	positive	1350777600	

	Summary	Text
0	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	"Delight" says it all	This is a confection that has been around a fe...
3	Cough Medicine	If you are looking for the secret ingredient i...
4	Great taffy	Great taffy at a great price. There was a wid...

```
In [55]: df.duplicated()
         df.head()
```

```
Out [55]:
```

	Id	ProductId	UserId	ProfileName	\
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	
2	3	B000LQOCHO	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham	"M. Wassir"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	1	1	5	1303862400	
1	0	0	1	1346976000	
2	1	1	4	1219017600	
3	3	3	2	1307923200	

4 0 0 5 1350777600

	Summary	Text
0	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	"Delight" says it all	This is a confection that has been around a fe...
3	Cough Medicine	If you are looking for the secret ingredient i...
4	Great taffy	Great taffy at a great price. There was a wid...

```
In [56]: sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True,)
```

```
In [57]: final=sorted_data.drop_duplicates(subset={'UserId', 'ProfileName', 'Time', 'Text'}
                                             ,keep='first',inplace=False)

final.shape
```

```
Out[57]: (393933, 10)
```

```
In [58]: (final["Id"].size*1.0)/(filtered_data["Id"].size*1.0)*100
```

```
Out[58]: 69.29901100176971
```

```
In [59]: display=pd.read_csv('Reviews.csv')
display
```

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [60]: print(final.shape)
```

```
(393931, 10)
```

```
In [61]: final['Score'].value_counts()
```

```
Out[61]: positive    336824
negative    57107
Name: Score, dtype: int64
```

```
In [62]: i=0;
for sent in final['Text'].values:
    if (len(re.findall('<.*?>', sent))):
        print(i)
        print(sent)
        break;
    i += 1;
```

0

In June
I saw a charming group
of roses all begin
to droop
I pepped them up

```

In [63]: import nltk
        from nltk.corpus import stopwords
        nltk.download('stopwords')
        stopwords = stopwords.words('english')
        import re
        import string
        from nltk.corpus import stopwords
        from nltk.stem import PorterStemmer
        from nltk.stem.wordnet import WordNetLemmatizer
        stop = set(stopwords.words('english')) #set of stopwords
        sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer
        def cleanhtml(sentence): #function to clean the word of any html-tags
            cleanr = re.compile('<.*?>')
            cleantext = re.sub(cleanr, ' ', sentence)
            return cleantext
        def cleanpunc(sentence): #function to clean the word of any punctuation or special_ch
            cleaned = re.sub(r'[?|!|\\'|"|#]',r'',sentence)
            cleaned = re.sub(r'[.,|)|(|\\|/]',r' ',cleaned)
            return cleaned
        print(stop)
        print('*****')
        print(sno.stem('tasty'))

```

```

[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\Saurabh\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
{'you', 'you'll', 'what', 'each', 'and', 'but', 'it's', 'shan', 'being', 'very', 'between', 'h
*****
tasti

```

```

In [64]: i=0
        str1=' '
        final_string=[]
        all_positive_words=[] # store words from +ve reviews here
        all_negative_words=[] # store words from -ve reviews here.
        s=''
        for sent in final['Text'].values:
            filtered_sentence=[]
            #print(sent);
            sent=cleanhtml(sent) # remove HTML tags
            for w in sent.split():
                for cleaned_words in cleanpunc(w).split():
                    if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                        if(cleaned_words.lower() not in stop):
                            s=(sno.stem(cleaned_words.lower())).encode('utf8')
                            filtered_sentence.append(s)
                            if (final['Score'].values)[i] == 'positive':

```

```

        all_positive_words.append(s) #list of all words us
    if(final['Score'].values)[i] == 'negative':
        all_negative_words.append(s) #list of all words used to descr
    else:
        continue
    else:
        continue
    #print(filtered_sentence)
    str1 = b" ".join(filtered_sentence) #final string of cleaned words
    #print("*****")

    final_string.append(str1)
    i+=1

```

In [65]: final['CleanedText']=final_string

2 Taking Sample Data

```

In [66]: n_samples = 25000
        df_sample = df.sample(n_samples)

        ###Sorting as we want according to time series

        df_sample.sort_values('Time',inplace=True)
        df_sample.head(10)

```

```

Out [66]:
      Id  ProductId  UserId  ProfileName \
451902  451903  B00004CXX9  A2DEE7F9XKP3ZR      jerome
230333  230334  B00004RYGX  A1GB1Q193DNFGR  Bruce Lee Pullen
451922  451923  B00004CXX9  ANIMV3SPDD8SH  Guy De Federicis
374333  374334  B00004CI84  A2HIZRVOKXKZ52      KAY N. FOWLER
374418  374419  B00004CI84  ADIDQRLLR4KBQ  "paradise_found"
374330  374331  B00004CI84  A10P3SQP78M1PP      James Gowen
230255  230256  B00004RYGX  AAI57M30XP5NK      "gibraud"
451937  451938  B00004CXX9  A1CAA94EOP0J2S  Travis J Smith
219437  219438  B00005IX97  A24IUOMVERXVEG  Rick "rick_street"
326756  326757  B000A28TJ2  A2RSOEBC1K70S      G. Preston

      HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
451902                    0                      1      5  959990400
230333                    5                      5      5  970531200
451922                    1                     12      2  992217600
374333                    0                      0      5  1012780800
374418                    2                      2      5  1015545600
374330                    0                      0      5  1015718400
230255                    0                      0      5  1025654400
451937                    4                      4      5  1036022400
219437                   11                     11      5  1036108800

```

326756 19 21 5 1036627200

```

Summary \
451902 Research
230333 Fabulous Comedic Fanasy Directed by a Master
451922 CASPER IS THE GHOST WITH THE MOST
374333 Beetlejuice - Great Fun for Everyone!
374418 Because My Middle Name is "Weird";
374330 A movie to "literally die for!!!"
230255 Love This Movie!
451937 Great Comedy
219437 Great Easy cup of Espresso
326756 Boxed Evolution

```

```

Text
451902 I'm getting crazy.<p>Is it really impossible t...
230333 Beetlejuice is an awe-inspiring wonderfully am...
451922 Michael Keaton brings no distinguishing charac...
374333 When vacationing Adam and Barbara Maitland mee...
374418 Set in a small New England town, Tim Burton's ...
374330 Beetlejuice is not a movie that you can't watc...
230255 This movie is a very odd movie but I love it b...
451937 Beetlejuice is the story of ghosts (Alec Baldw...
219437 I've purchased both the Espressione Espresso (...
326756 Long live Camper Van Beethoven!! This is a ble...

```

```

In [67]: new_final_counts=final[:25000]
         print(new_final_counts.shape)

```

(25000, 11)

3 An Introduction to Bag Of Words =>

```

In [ ]: We convert text to a numerical representation called a feature vector.
        A feature vector can be as simple as a list of numbers.

```

The bag-of-words model is one of the feature extraction algorithms for text.

4 Classification Of Bag Of Words Using Knn (Brute Force)

```

In [ ]: from sklearn.model_selection import train_test_split

        from sklearn import preprocessing
        X_train, X_test, y_train, y_test = train_test_split(new_final_counts
        ['CleanedText'].values,new_final_counts['Score'].values ,test_size=0.30,
        shuffle=False)

```



```

In [ ]: #Text -> Uni gram Vectors

uni_gram = CountVectorizer()
X_train = uni_gram.fit_transform(X_train)

#Normalize Data

X_train = preprocessing.normalize(X_train)
print("Train Data Size: ",X_train.shape)
X_test = uni_gram.transform(X_test)

#Normalize Data

X_test = preprocessing.normalize(X_test)
print("Test Data Size: ",X_test.shape)

In [ ]: from sklearn.model_selection import TimeSeriesSplit

tscv = TimeSeriesSplit(n_splits=10)
for train, cv in tscv.split(X_train):

    # print("%s %s" % (train, cv))
    print(X_train[train].shape, X_train[cv].shape)

In [ ]: %time
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(algorithm='brute')

# neigh = np.arange(1,100,2)

param_grid = {'n_neighbors':np.arange(1,100,2)} #params we need to try on classifier
tscv = TimeSeriesSplit(n_splits=10) #For time based splitting

gsv = GridSearchCV(knn,param_grid,cv=tscv,verbose=1)
gsv.fit(X_train,y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))

In [69]: #Testing Accuracy on Test data

%time

from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix

```

```

knn = KNeighborsClassifier(n_neighbors=11)
knn.fit(X_train,y_train)
y_pred = knn.predict(X_test)

print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))
print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")

df_cm = pd.DataFrame(confusion_matrix(y_test, y_pred), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')

```

Wall time: 0 ns


Accuracy on test set: 85.973%

Confusion Matrix of test set:

[[TN FP]

[FN TP]]

Out[69]: <matplotlib.axes._subplots.AxesSubplot at 0x205ff6be2b0>



5 Observation

In []: 1. Here we observe that while using 25k points the heat map of confusion matrix is very clear and a/c to this we see that the model is perfect

because the TN and TP are having high values.

2.The accuracy on text data using brute force algorithm is 85.973% which shows that model is perfect.

6 An introduction to TF-IDF =>

In []: TF-IDF stands for Term Frequency|Inverse Data Frequency|.

Term Frequency (tf): gives us the frequency of the word in each document in the corpus. It is the ratio of number of times the word appears in a document compared to the total number of words in that document. It increases as the number of occurrences of that word within the document increases. Each document has its own tf.

Inverse Data Frequency (idf)=>

Used to calculate the weight of rare words across all documents in the corpus. The words that occur rarely in the corpus have a high IDF score. It is given by the equation below.

7 Classification of TF-IDF using Knn Algorithm (Brute Force)

```
In [81]: %%time
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn import preprocessing

#Breaking into Train and test
X_train, X_test, y_train1, y_test = train_test_split(new_final_counts
['CleanedText'].values,new_final_counts['Score'].values ,test
_size=0.30,shuffle=False)

tfidf = TfidfVectorizer(ngram_range=(1,2)) #Using bi-grams
X_train_tfidf=tfidf.fit_transform(X_train)

#Normalize Data
X_train_tfidf_counts = preprocessing.normalize(X_train_tfidf)
print("Train Data Size: ",X_train_tfidf_counts.shape)
X_test = tfidf.transform(X_test)
#Normalize Data
X_test = preprocessing.normalize(X_test)
print("Test Data Size: ",X_test.shape)
```

Train Data Size: (17500, 379637)

Test Data Size: (7500, 379637)

Wall time: 3.55 s

In []: %time

```
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(algorithm='brute')

# neigh = np.arange(1,100,2)

param_grid = {'n_neighbors':np.arange(1,100,2)} #params we need to try on classifier
tscv = TimeSeriesSplit(n_splits=10) #For time based splitting

gsv = GridSearchCV(knn,param_grid,cv=tscv,verbose=1)
gsv.fit(X_train,y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

In [70]: #Testing Accuracy on Test data

```
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=9)
knn.fit(X_train,y_train)
y_pred = knn.predict(X_test)

print("Accuracy on test set: %.3f%%"%(accuracy_score(y_test, y_pred)*100))
print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")


df_cm = pd.DataFrame(confusion_matrix(y_test, y_pred), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')
```

Accuracy on test set: 86.107%

Confusion Matrix of test set:

```
[ [TN FP]
  [FN TP] ]
```

Out[70]: <matplotlib.axes._subplots.AxesSubplot at 0x205864517f0>



8 Observation=>

- In []: 1. Here we observe that while using 25k points the heat map of confusion matrix is very clear and a/c to this we see that the model is perfect because the TN and TP are having high values.
2. The accuracy on text data using brute force algorithm is 86.107% which shows that model is perfect.

9 Word2vec =>

In []: word2vec is an algorithm for constructing vector representations of words, also known as word embeddings. The vector for each word is a semantic description of how that word is used in context, so two words that are used similarly in text will get similar vector representations. Once you map words into vector space, you can then use vector math to find words that have similar semantics.

10 Using Google News Word2Vectors

In []: Gensim

Gensim is a robust open-source vector space modeling and topic modeling toolkit implemented in Python. It uses NumPy, SciPy and optionally Cython for performance. Gensim is specifically designed to handle

large text collections, using datastreaming and efficient incremental algorithms, which differentiates it from most other scientific software packages that only target batch and in-memory processing.

```
In [ ]: %time
```

```
import re
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

model = KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin',
                                          binary=True)
```

```
In [83]: # Train your own Word2Vec model using your own text corpus
```

```
import gensim
i=0
list_of_sent=[]
for sent in final['Text'].values:
    filtered_sentence=[]
    sent=cleanhtml(sent)
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower())
            else:
                continue
    list_of_sent.append(filtered_sentence)
```

```
In [84]: w2v_model=gensim.models.Word2Vec(list_of_sent,min_count=5,size=50, workers=4)
```

```
In [85]: w2v_vocub = model.wv.vocab
len(w2v_vocub)
```

```
Out[85]: 3000000
```

11 Avg Word2Vec Using Brute Force Algorithm

```
In [ ]: 1.One of the most naive but good ways to convert a sentence into a vector.
        2.Convert all the words to vectors and then just take the avg of the
        vectors the resulting vector represent the sentence.
```

```
In [86]: # average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
```

```

    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))

```

393931

```

In [ ]: %time
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(algorithm='brute')

# neigh = np.arange(1,100,2)

param_grid = {'n_neighbors':np.arange(1,100,2)} #params we need to try on classifier
tscv = TimeSeriesSplit(n_splits=10) #For time based splitting

gsv = GridSearchCV(knn,param_grid,cv=tscv,verbose=1)
gsv.fit(X_train,y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))

```

In [71]: #Testing Accuracy on Test data

```

from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=11)
knn.fit(X_train,y_train)

y_pred = knn.predict(X_test)

print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))
print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")

df_cm = pd.DataFrame(confusion_matrix(y_test, y_pred), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')

```


Accuracy on test set: 85.973%

Confusion Matrix of test set:

```
[ [TN FP]
```

[FN TP]]

Out[71]: <matplotlib.axes._subplots.AxesSubplot at 0x20581de7d68>



12 Observation

- In []: 1. Here we observe that while using 25k points the heatmap of confusion matrix is very clear and a/c to this we see that the model is perfect because the TN and TP are having high values.
2. The accuracy on text data using brute force algorithm is 85.973% which shows that model is perfect.

13 Tf-idf W2Vec Using Brute Force Algorithm

- In []: 1. Another way to convert sentence into vectors
2. Take weighted sum of the vectors divided by the sum of all the tfidf's i.e. $(\text{tfidf}(\text{word}) \times \text{w2v}(\text{word})) / \text{sum}(\text{tfidf's})$

In [87]: *# TF-IDF weighted Word2Vec*
tfidf_feat = tfidf.get_feature_names() *# tfidf words/col-names*
final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf


```

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this l
row=0;
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tfidf = final_tf_idf[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        except:
            pass
    sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1

```

```

In [88]: new_tfidf_sent_vectors=tfidf_sent_vectors[:25000]
        new_final=final[:25000]

```

```

In [90]: #Not shuffling the data as we want it on time basis
        X_train, X_test, y_train, y_test = train_test_split(new_tfidf_sent
        _vectors,new_final['Score'].values,test_size=0.3,shuffle=False)

```

```

In [91]: from sklearn.model_selection import TimeSeriesSplit
        tscv=TimeSeriesSplit(n_splits=10)

```

```

In [93]: from sklearn.model_selection import TimeSeriesSplit
        tscv=TimeSeriesSplit(n_splits=10)

```

```

In [95]: X_train=np.isnan(X_train)
        np.where(np.isnan(X_train))
        np.nan_to_num(X_train)
        X_test=np.isnan(X_test)
        np.where(np.isnan(X_test))
        np.nan_to_num(X_test)

```

```

Out[95]: array([[ True,  True,  True, ...,  True,  True,  True],
               [ True,  True,  True, ...,  True,  True,  True],
               [ True,  True,  True, ...,  True,  True,  True],
               ...,
               [ True,  True,  True, ...,  True,  True,  True],
               [ True,  True,  True, ...,  True,  True,  True],
               [ True,  True,  True, ...,  True,  True,  True]])

```

```

In [96]: %time

```

```

from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import TimeSeriesSplit

knn = KNeighborsClassifier(algorithm='brute')

# neigh = np.arange(1,100,2)

param_grid = {'n_neighbors':np.arange(1,40,2)} #params we need to try on classifier
tscv = TimeSeriesSplit(n_splits=10) #For time based splitting

gsv = GridSearchCV(knn,param_grid,cv=tscv,verbose=1)
gsv.fit(X_train,y_train)
print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))

```

Wall time: 0 ns

Fitting 10 folds for each of 20 candidates, totalling 200 fits

Best HyperParameter: {'n_neighbors': 1}

Best Accuracy: 85.67%

[Parallel(n_jobs=1)]: Done 200 out of 200 | elapsed: 8.0min finished

In [97]: *#Testing Accuracy on Test data*

```

from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=9)
knn.fit(X_train,y_train)
y_pred = knn.predict(X_test)

print("Accuracy on test set: %.3f%%"%(accuracy_score(y_test, y_pred)*100))
print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")

df_cm = pd.DataFrame(confusion_matrix(y_test, y_pred), range(2),range(2))
sns.set(font_scale=1.4)#for label size
sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g')

```

Accuracy on test set: 85.427%


Confusion Matrix of test set:

```

[ [TN FP]
  [FN TP] ]

```

Out[97]: <matplotlib.axes._subplots.AxesSubplot at 0x205fdd9c860>



14 Conclusions

In []: Here is the test accuracies of Feature Vectors

```
Bow          = 85.973%
TF-IDF       = 86.107%
Avg word2vec = 85.973%
TF-IDF W2V   = 85.427%
```

In []: Note: As I have taken only 25k points(due to huge training time) the accuracy will not be the representative of the real accuracy

1. Best Accuracy of 85.107% is achieved by TF-IDF Featurization
2. The kd-tree and brute implementation of KNN gives relatively similar results. But in this I'am not using kdtree because when I pass the sparse matrix to knn with kdtree algo, it can't process it, so it will internally stop using kdtree and uses brute force only
3. KNN is a very slow Algorithm compared to others takes a lot of time to train good for this dataset